

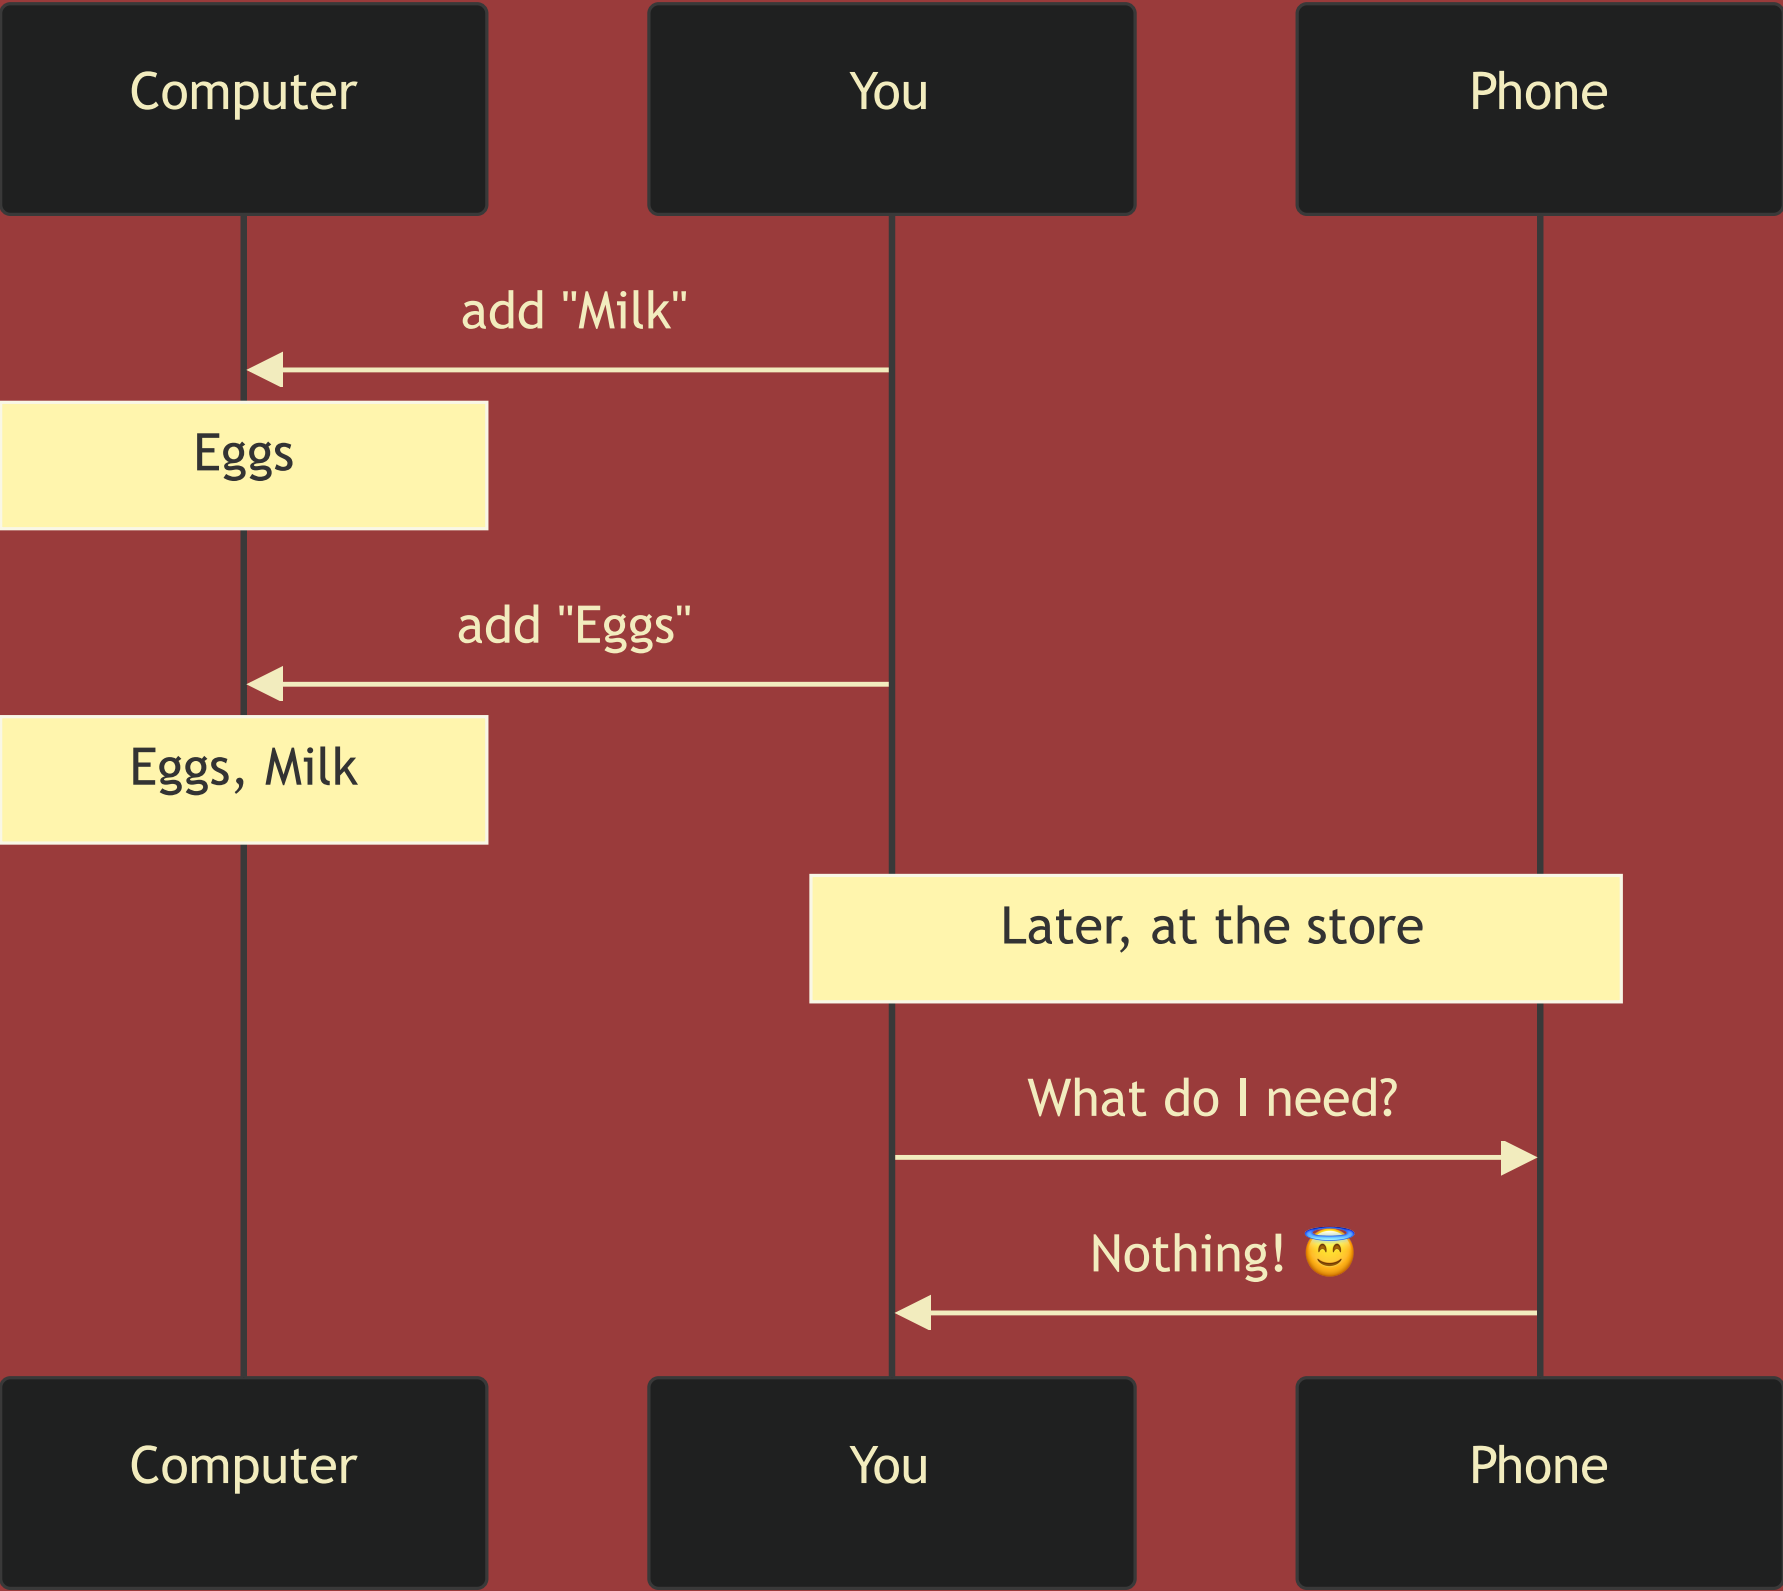
CRDTs in Rust

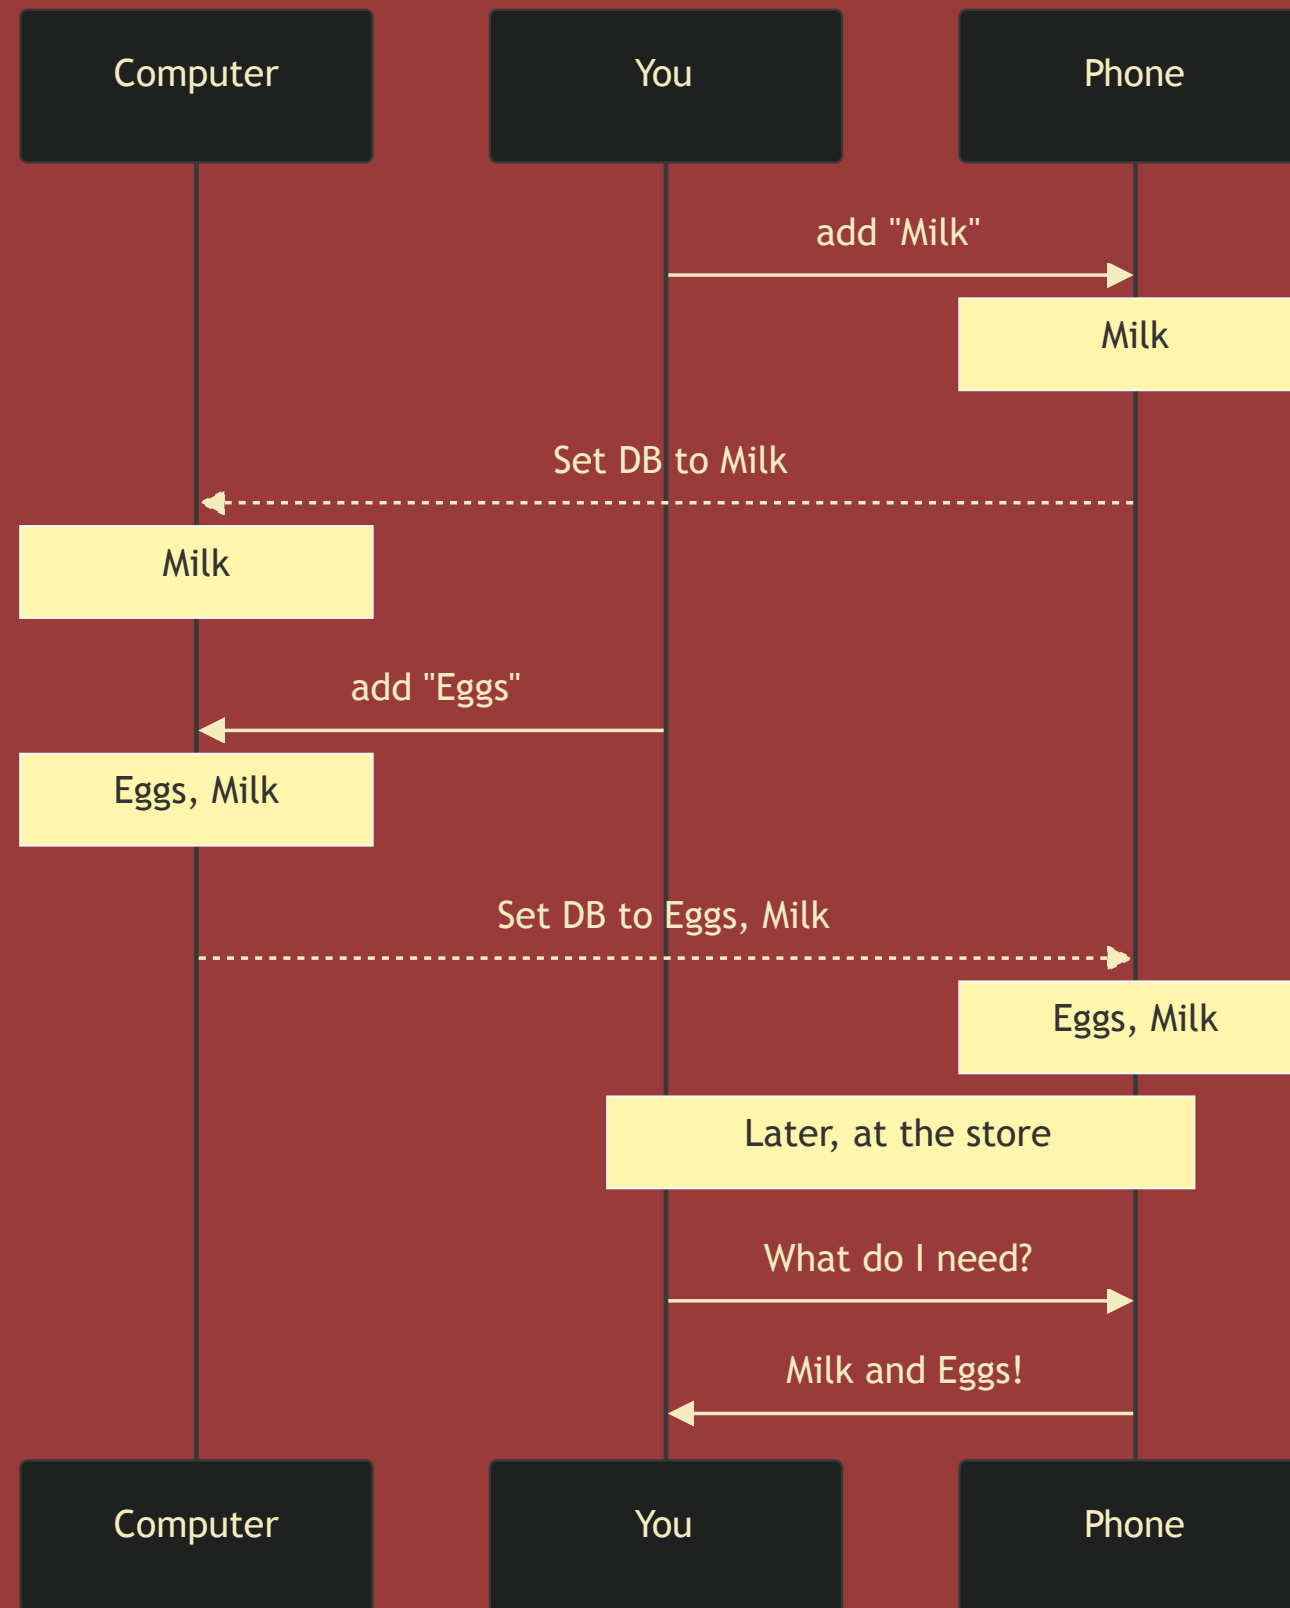


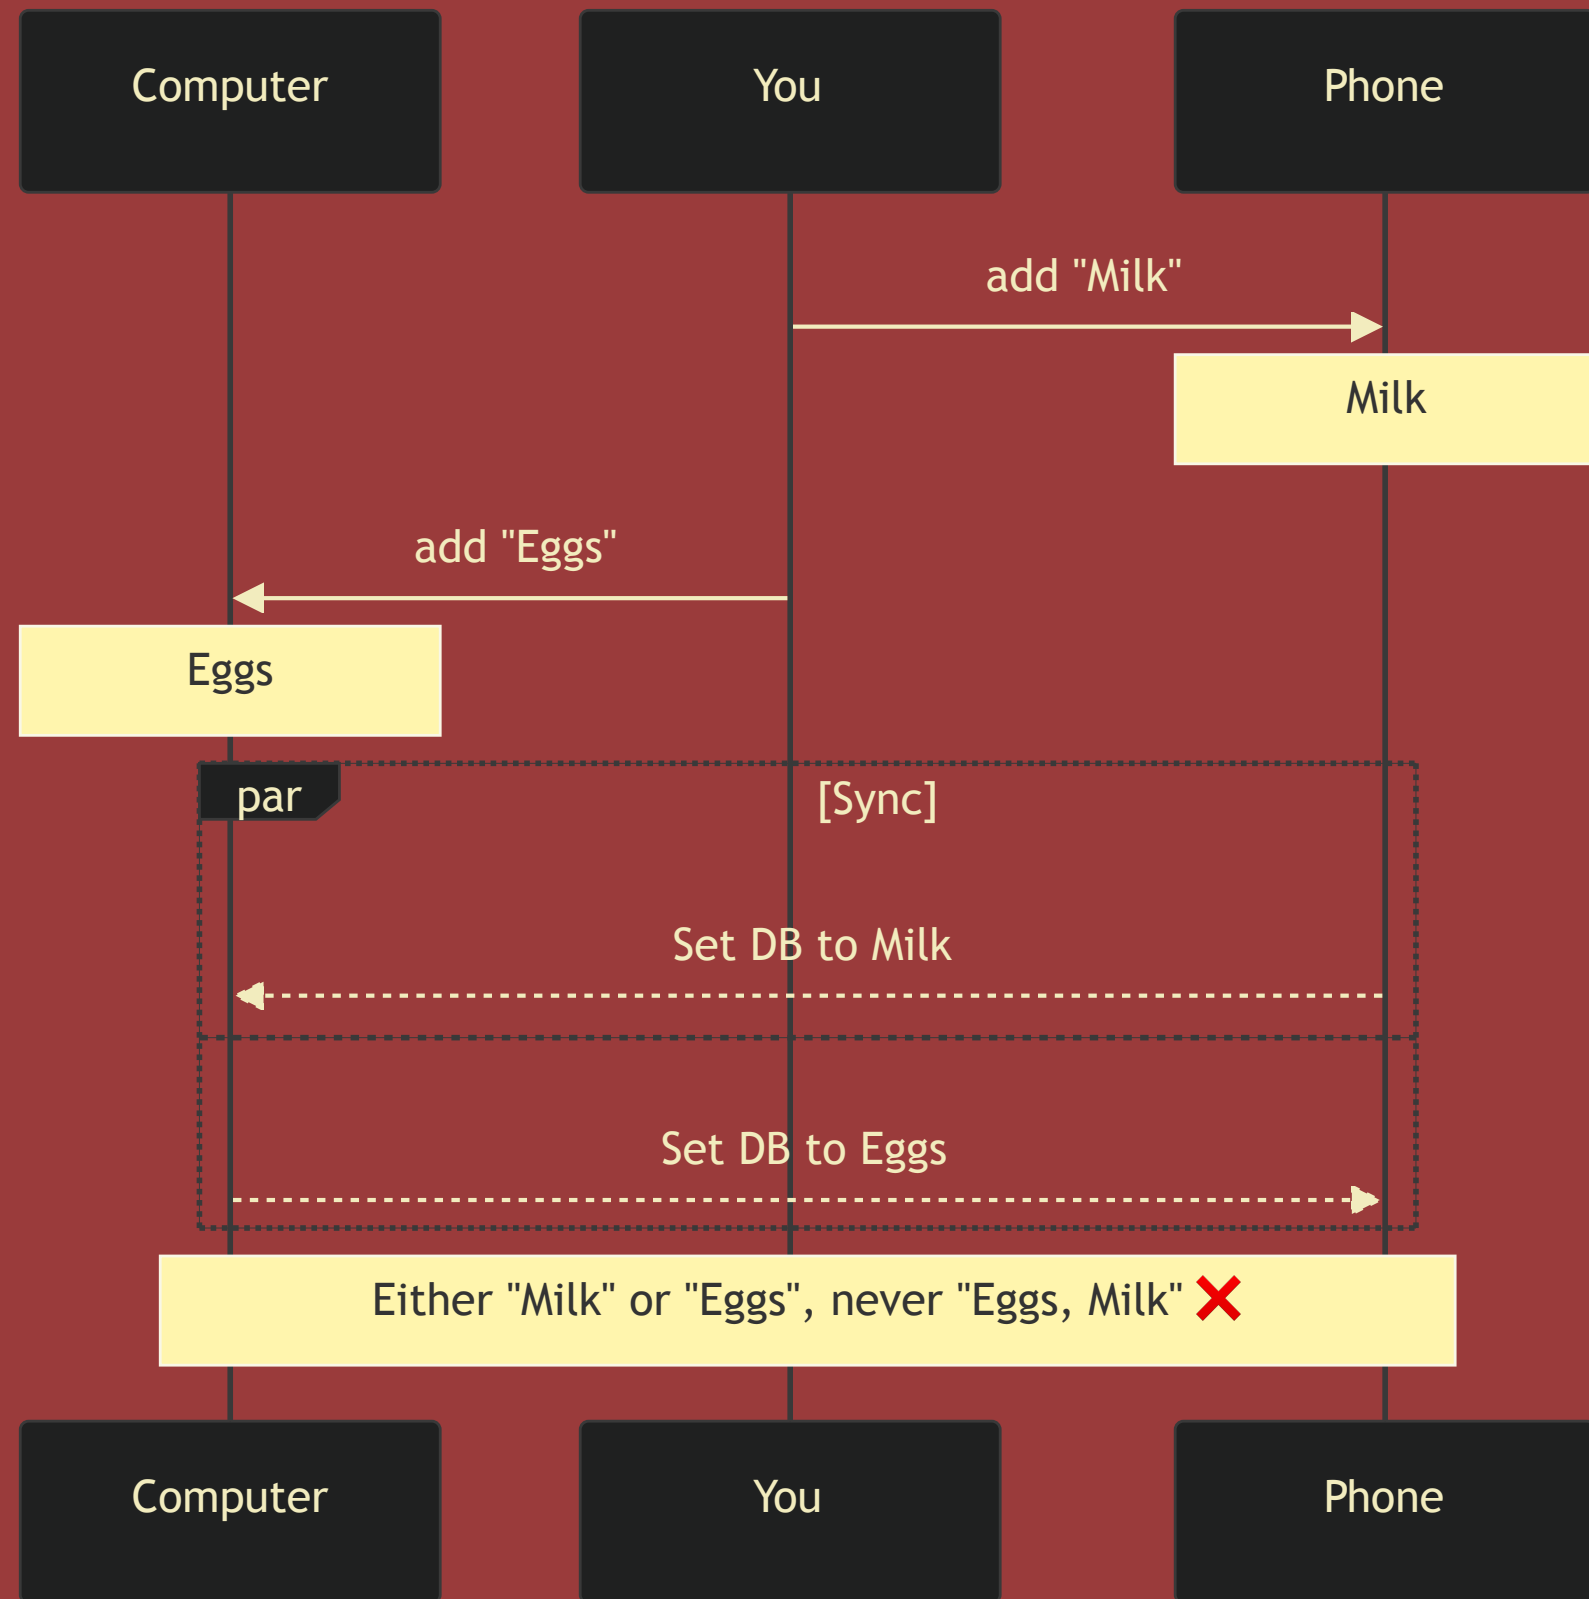
He11o!

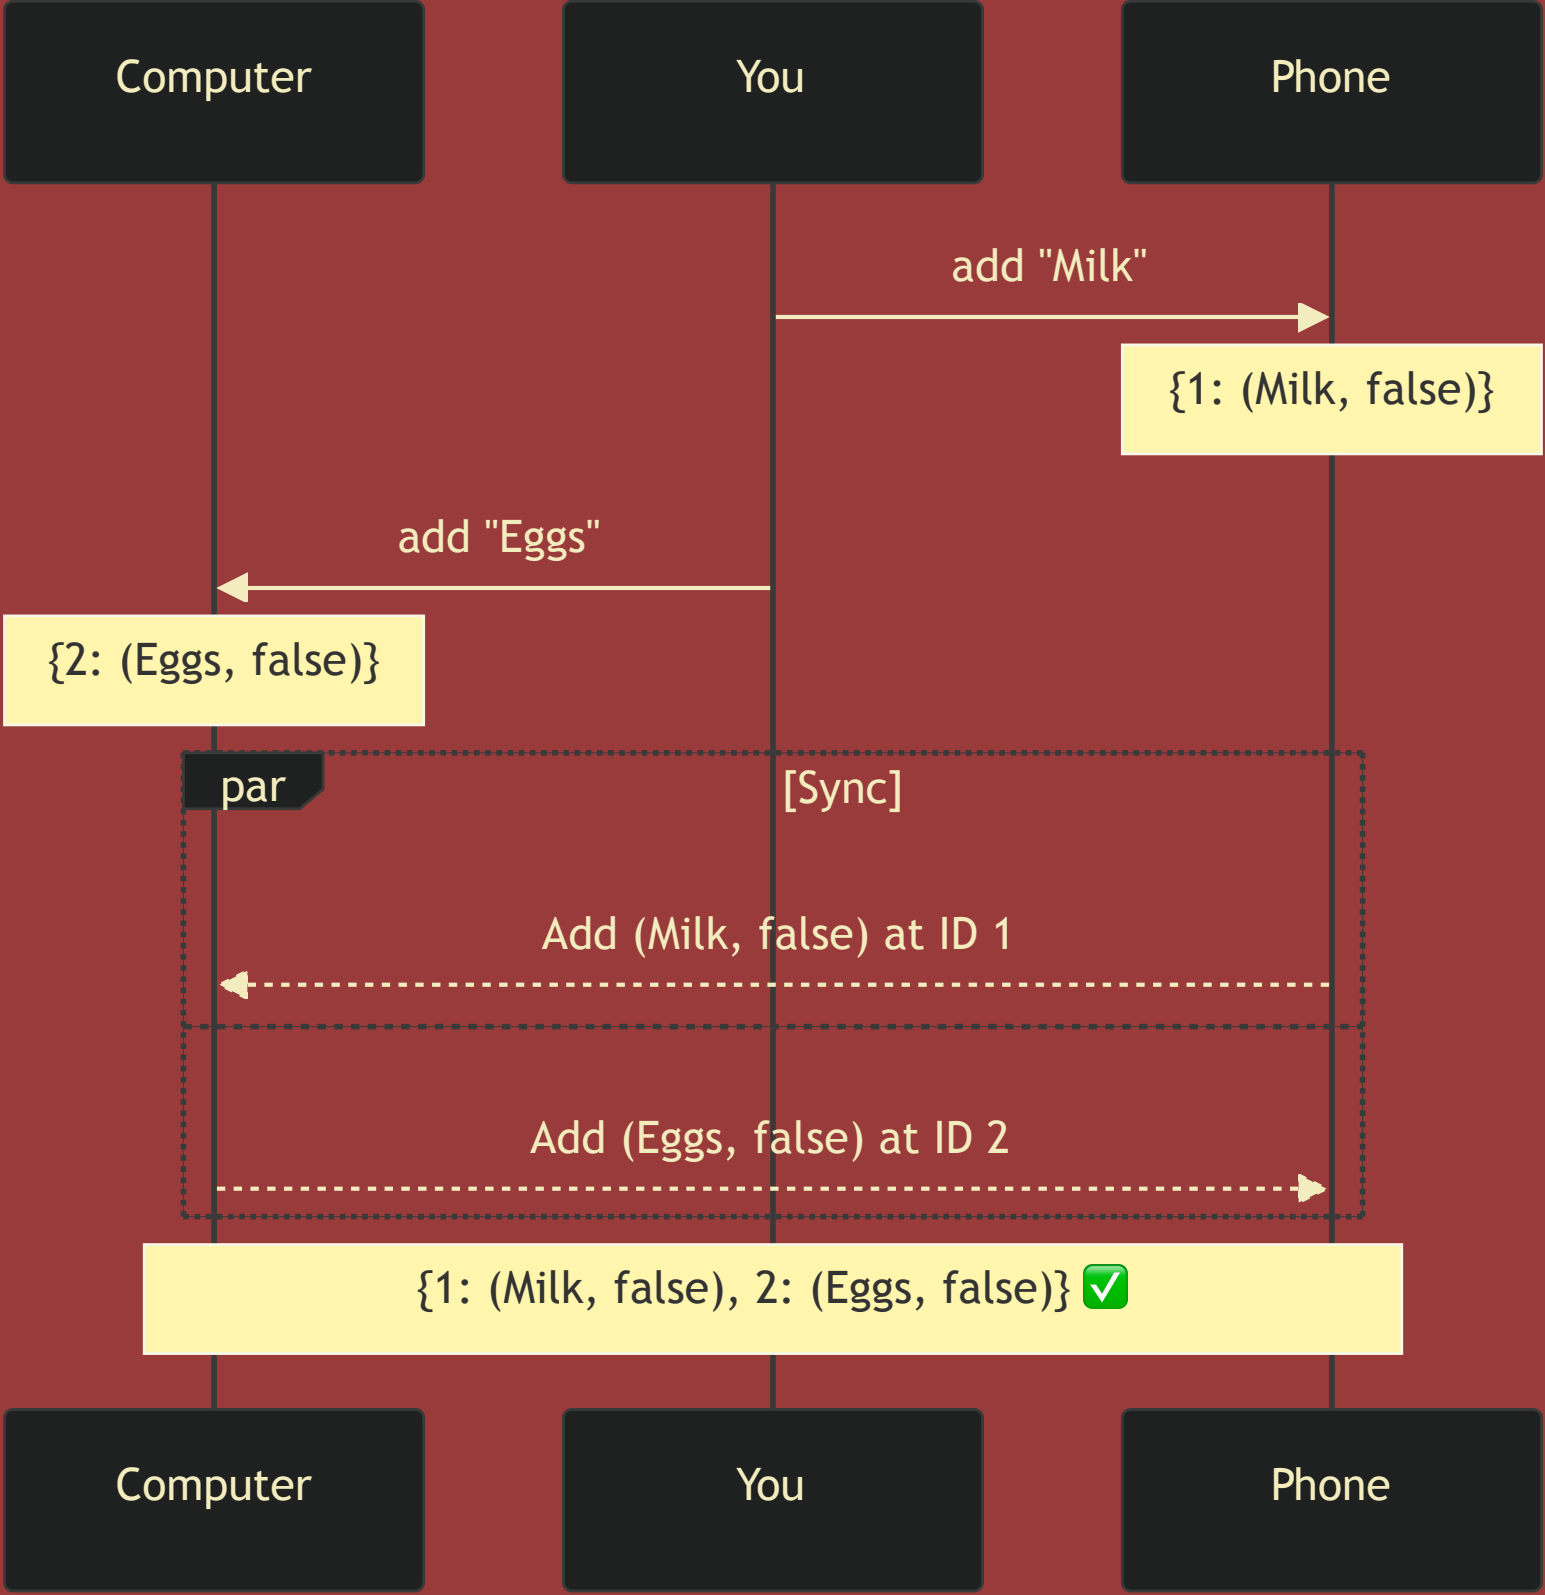
The Situation: to-do list

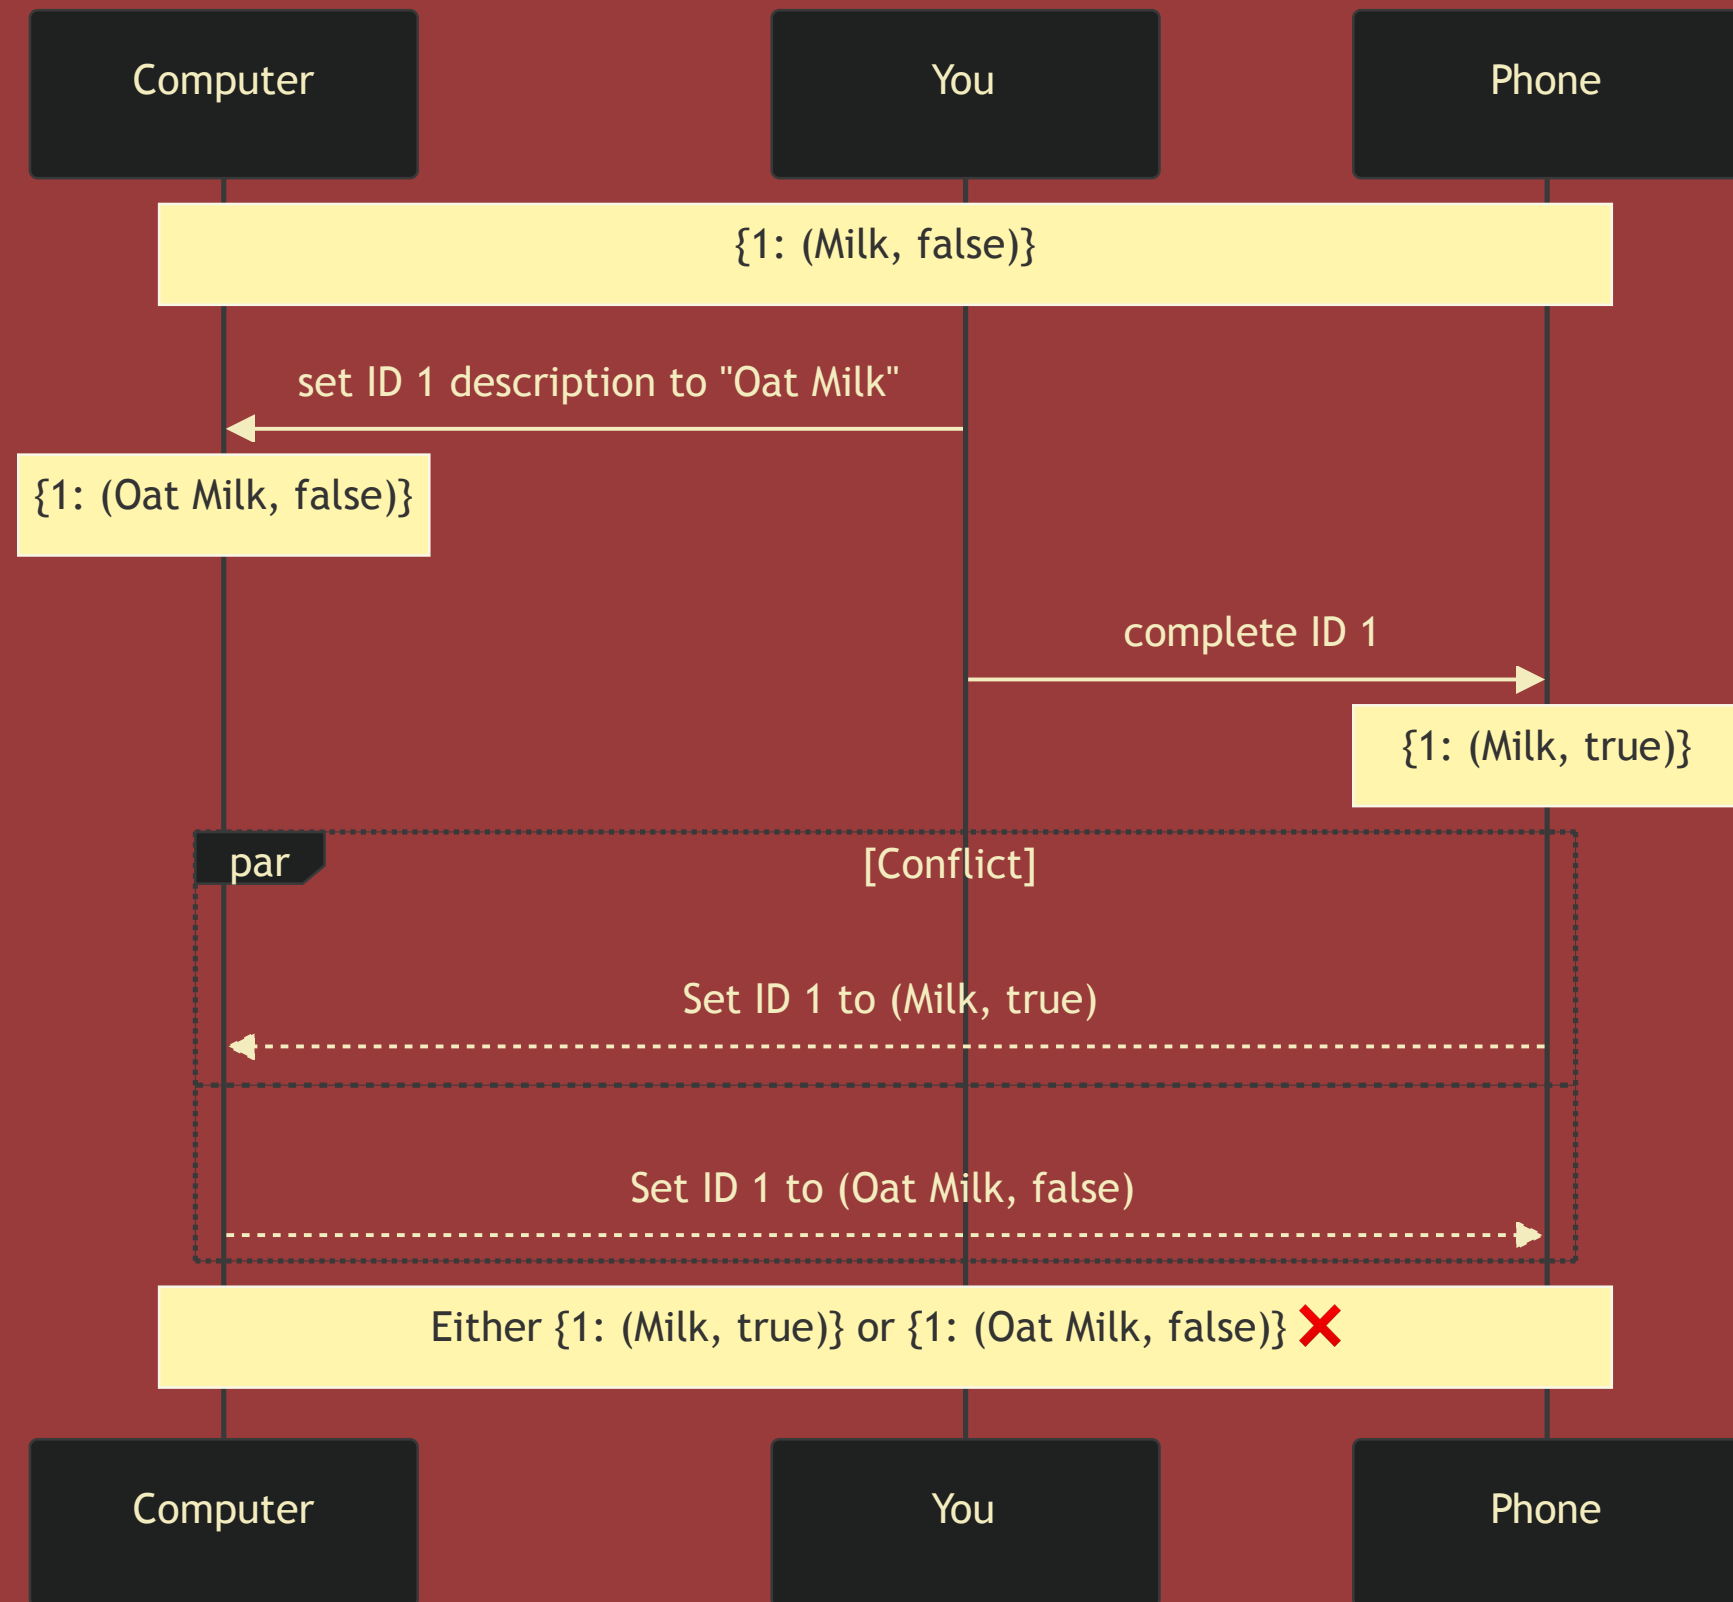


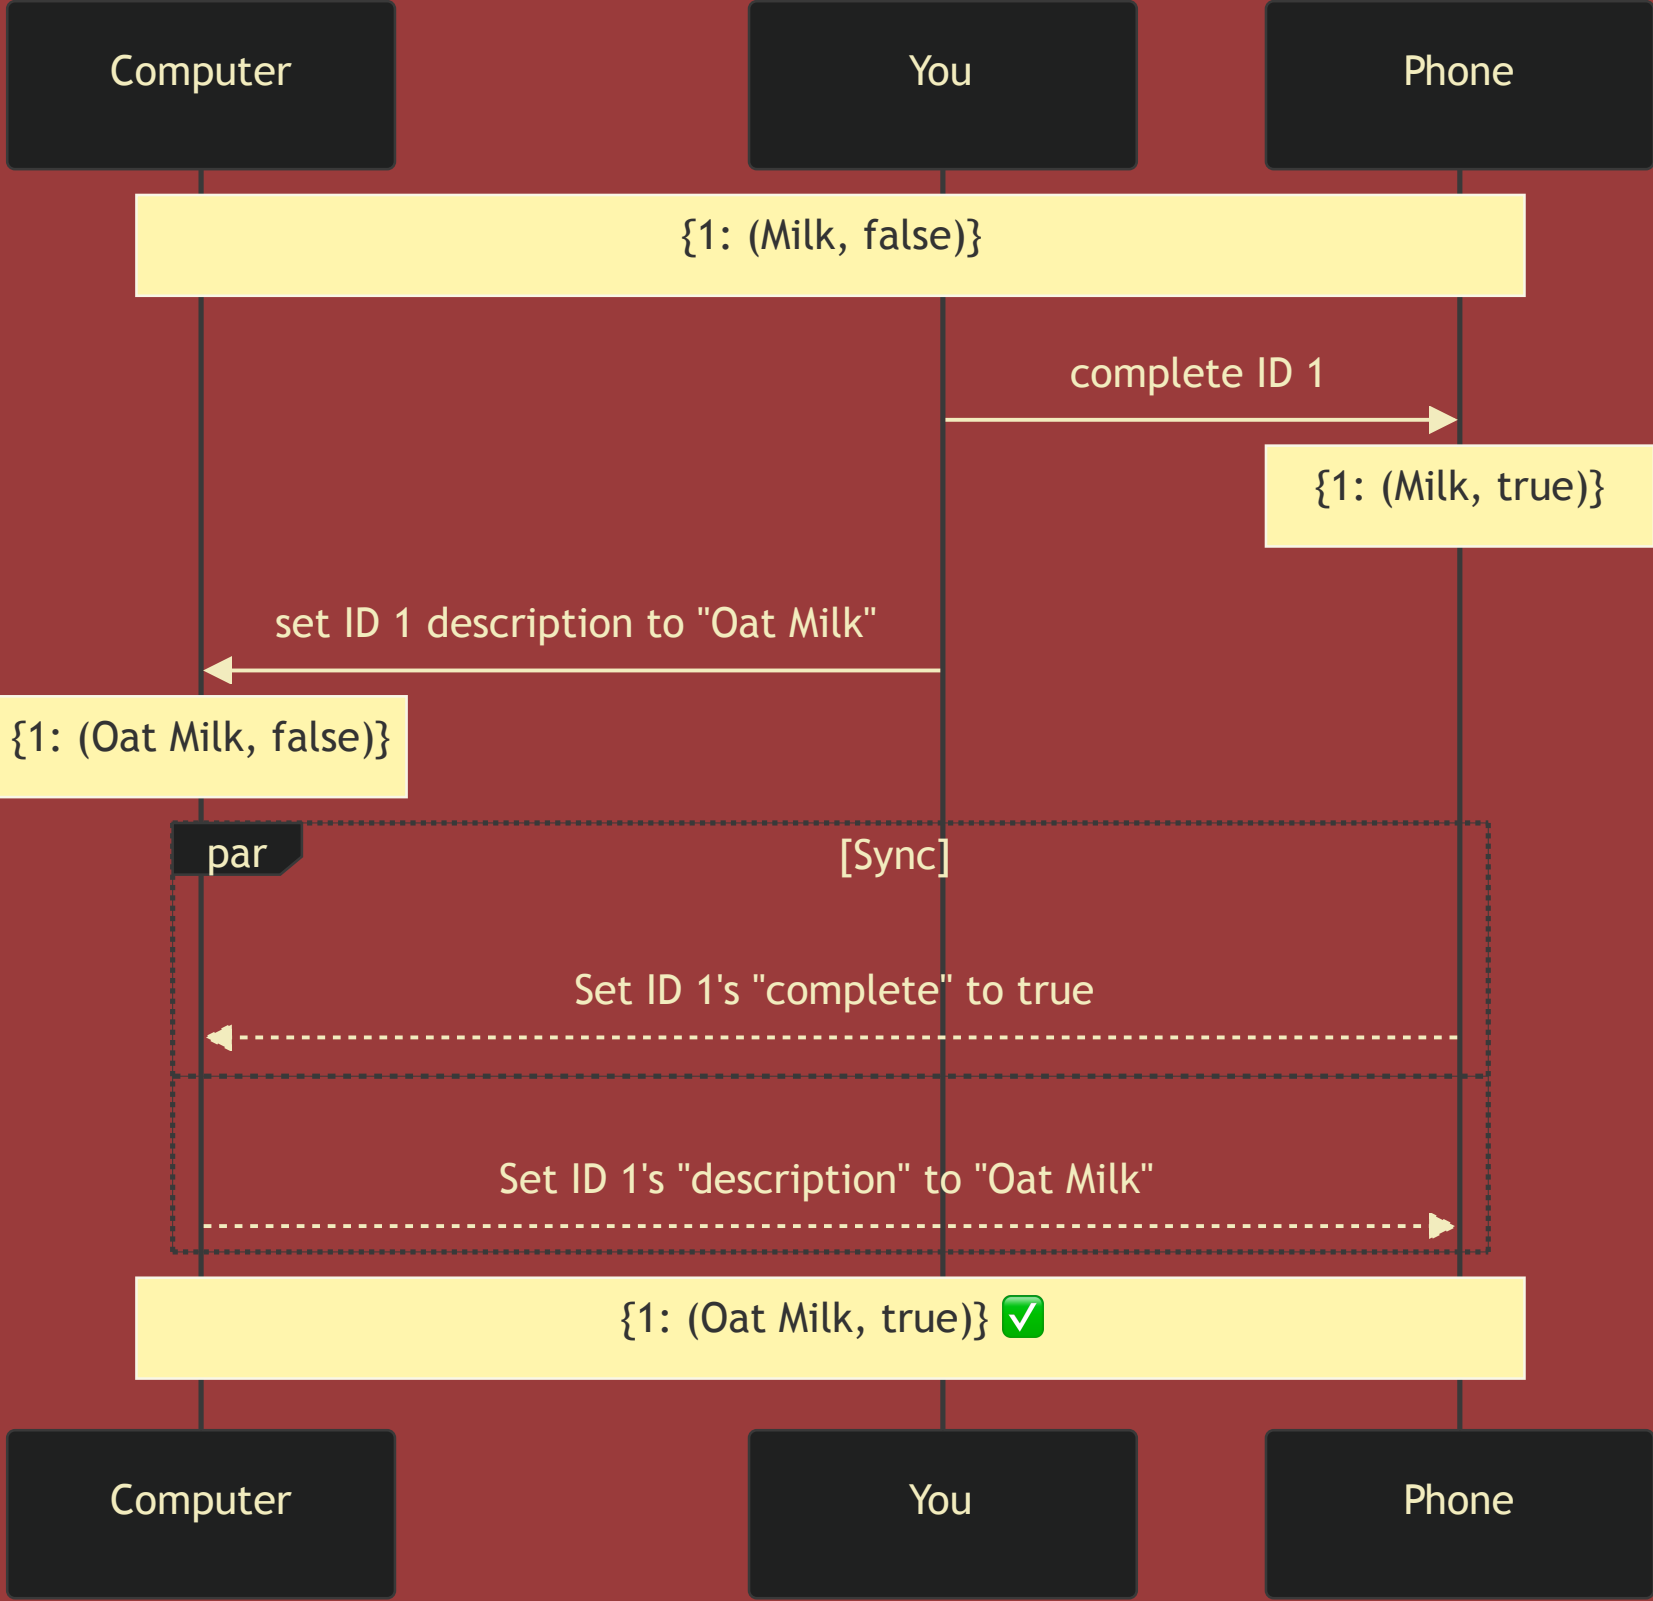


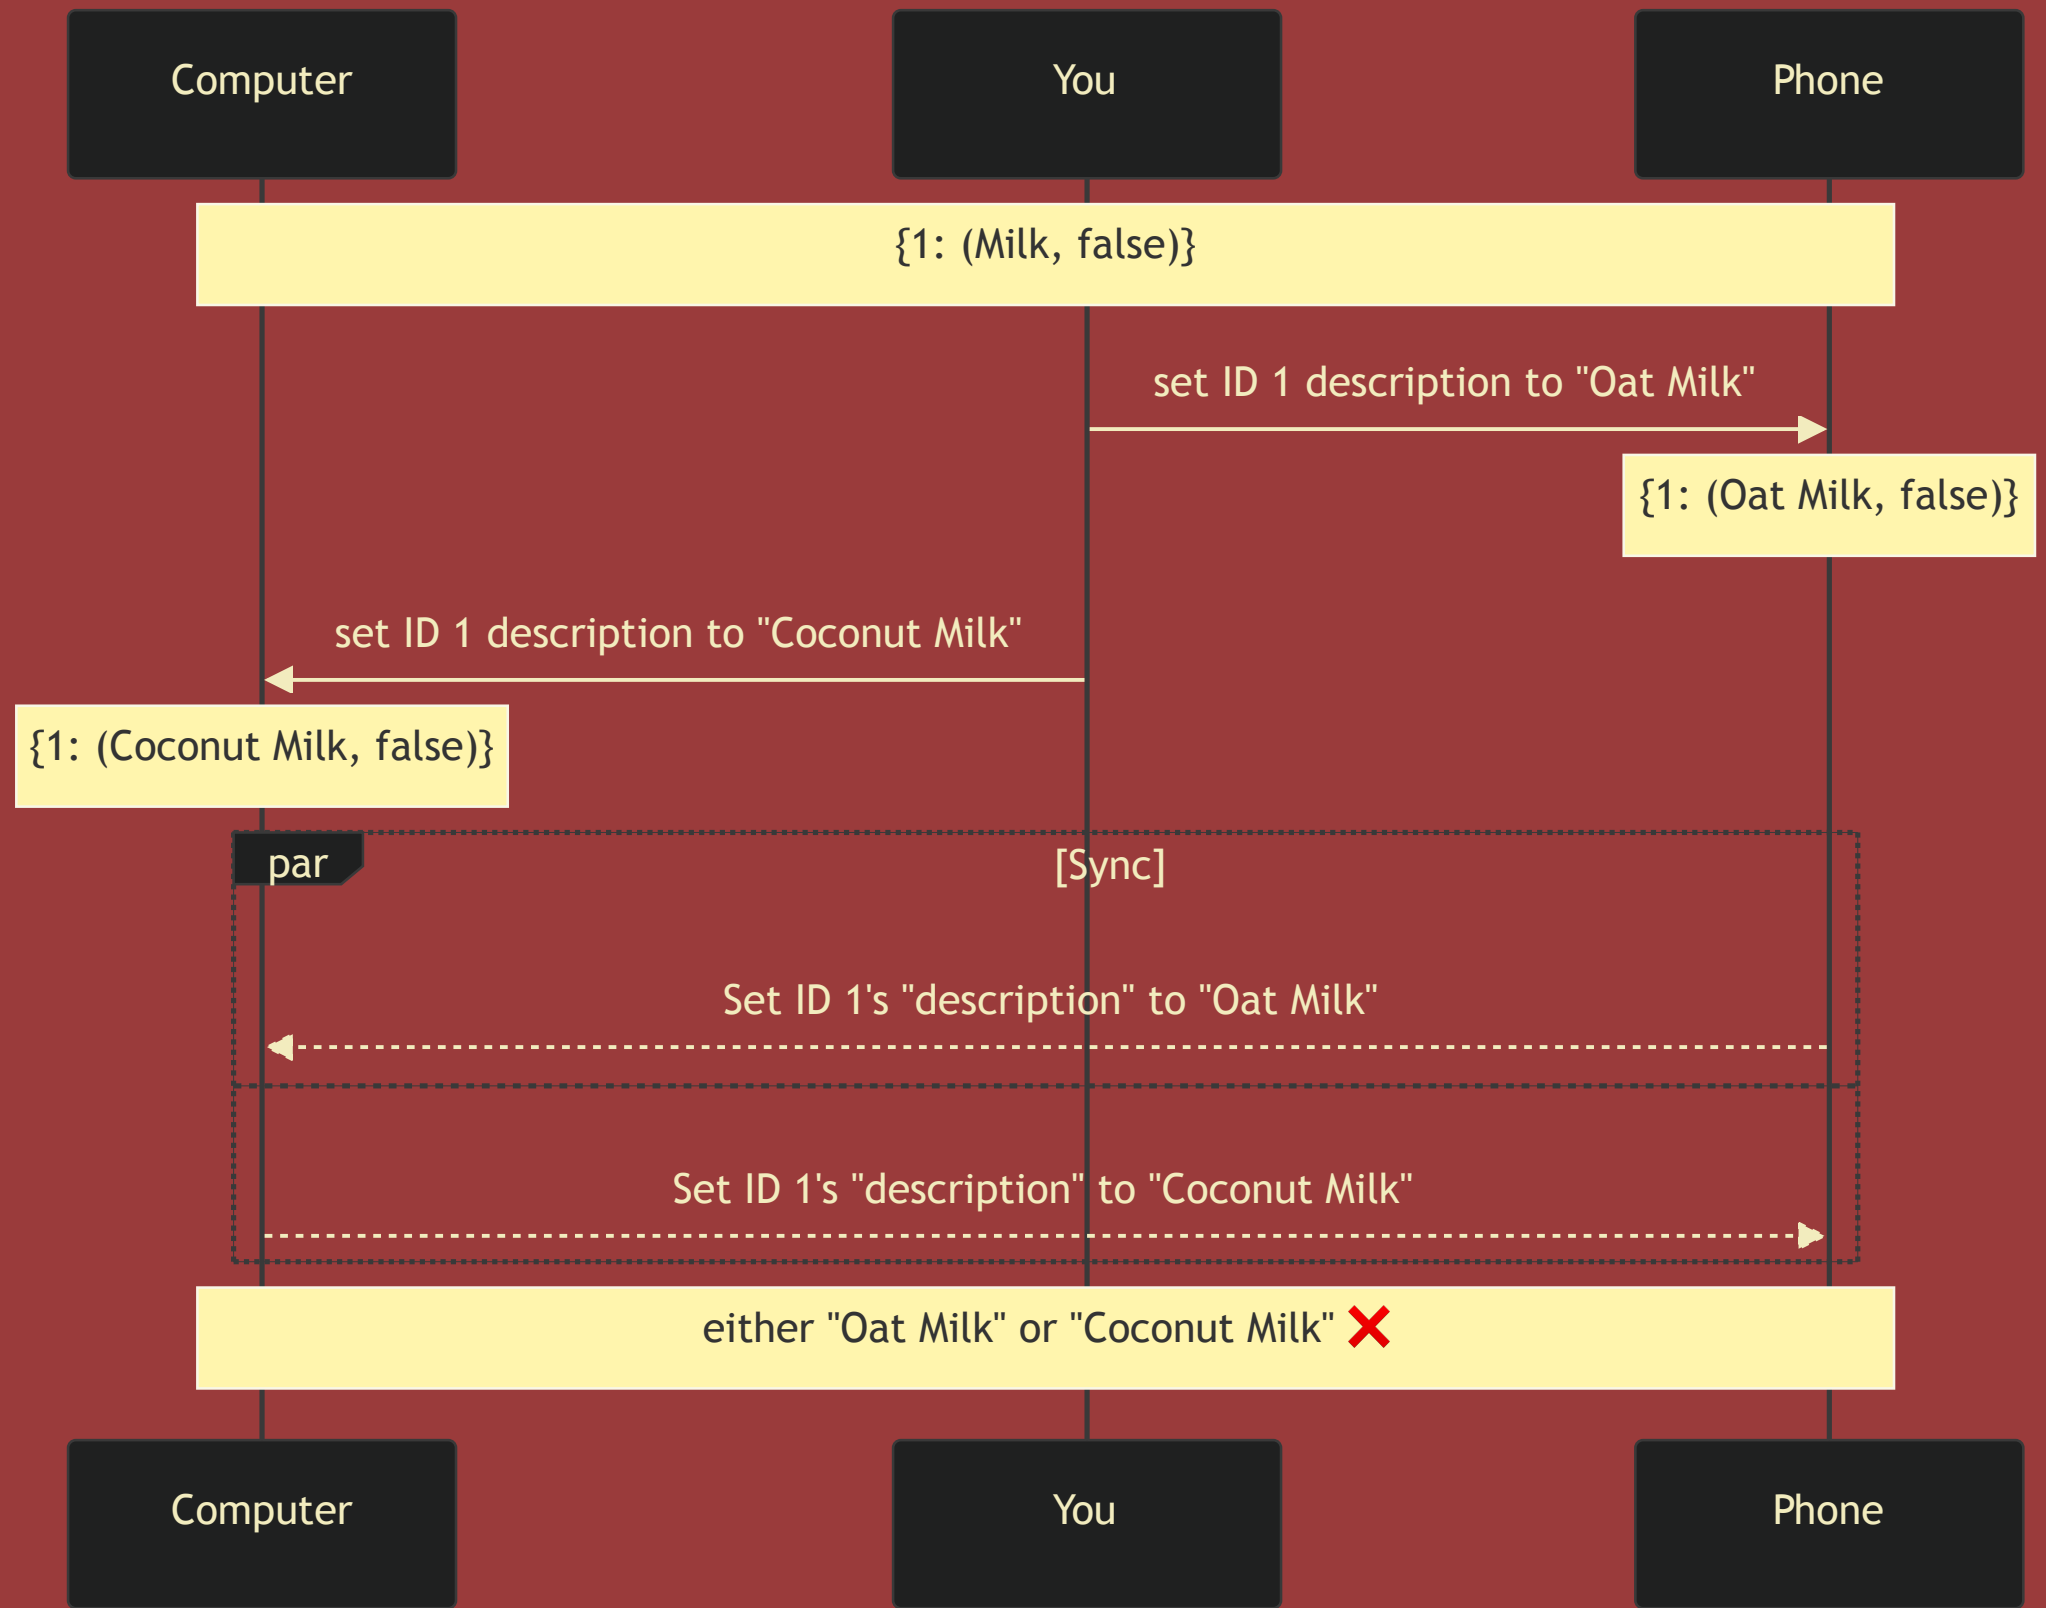












The Promise of CRDTs:

Merge any two values consistently, no matter when or where they were created.

But first of all... what are they?

- Conflict-Free (sometimes Convergent) Replicated Data Types
- Collaboration: Yjs, Apple Notes
- Databases: Redis, Riak, Antidote
- Local-first apps

Making a to-do list

We need to:

1. Make a state-based CRDT that can...
 - Keep a set of tasks by ID
 - Resolve concurrent edits in any fields
 - Eventually archive (delete) completed tasks
2. Decide on a sync protocol
3. ✨ Demo ✨

Building a (state-based) CRDT

You need two things:

1. A data structure (your state)
2. A merge function, which is:
 - Idempotent ($a \bullet a = a$) so merging identical replicas is OK.
 - Commutative ($a \bullet b = b \bullet a$) so merging in any direction is OK.
 - Associative ($a \bullet (b \bullet c) = (a \bullet b) \bullet c$) so merging in any order is OK.

Merge

```
pub trait Merge {  
    fn merge(self, other: Self) -> Self;  
}
```


Merge

```
pub trait Merge {  
    fn merge_mut(&mut self, other: Self);  
}
```

The simplest CRDT

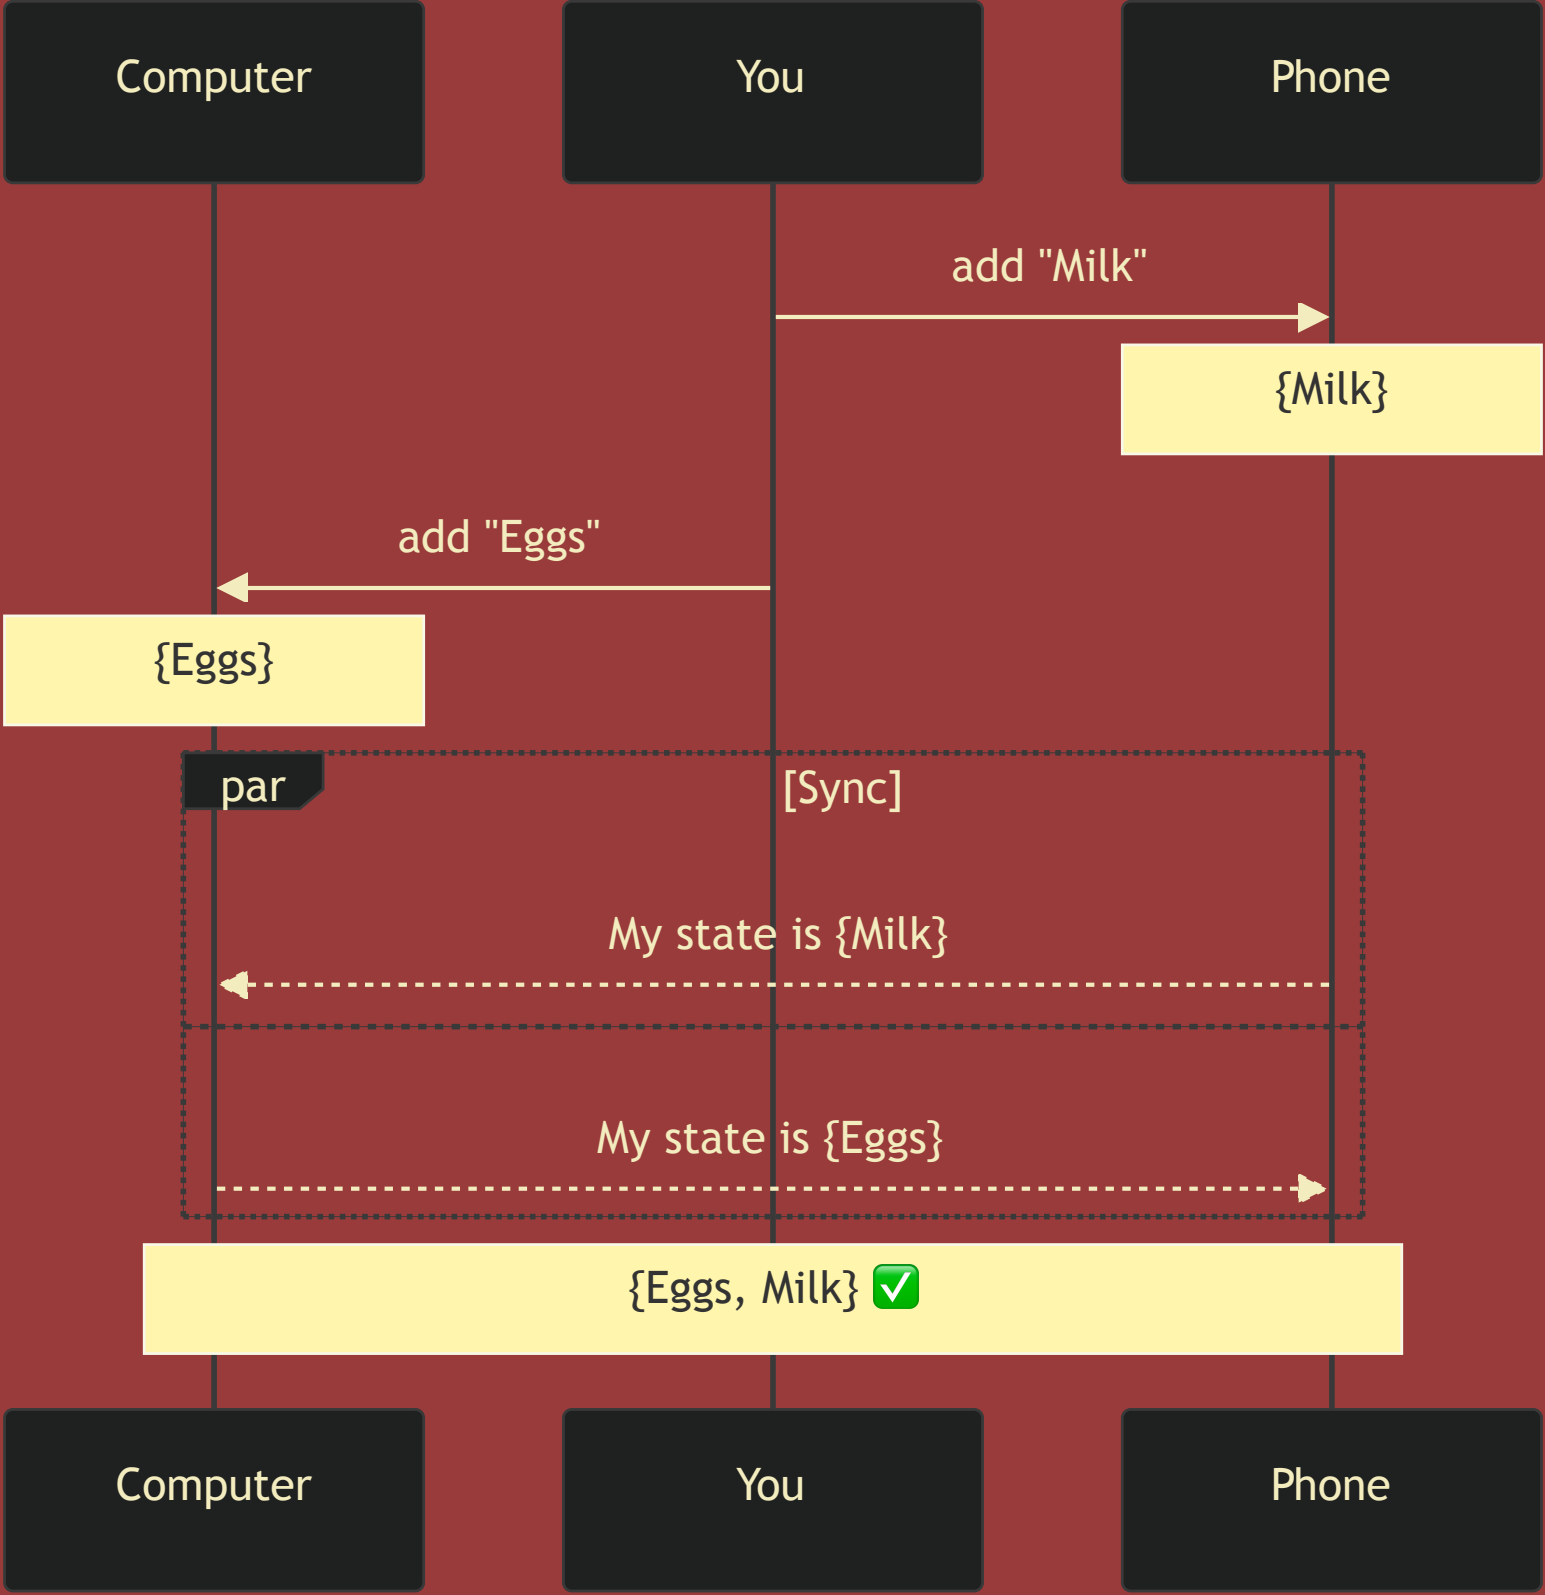
The "grow-only set", or G-Set. The merge function is set union.

The simplest CRDT

The "grow-only set", or G-Set. The merge function is set union.

```
pub struct GSet<T: Eq + Ord>(BTreeSet<T>);

impl<T: Eq + Ord> Merge for GSet<T> {
    fn merge_mut(&mut self, mut other: Self) {
        self.0.append(&mut other.0);
    }
}
```



Sync Protocol

1. Send your whole state to every peer.
2. When you get a state, merge it with yours.
 - Duplicate/redundant data is a problem
 - Delta-state CRDTs help with this

Keeping Fields

The "last write wins register", or LWW-Register. Keep any value, merge by taking the last written value.

- ... but what does "last write" mean?
- We can't use wall time because clocks drift.
- We could use a logical clock but having physical time is nice.

Hybrid Logical Clock

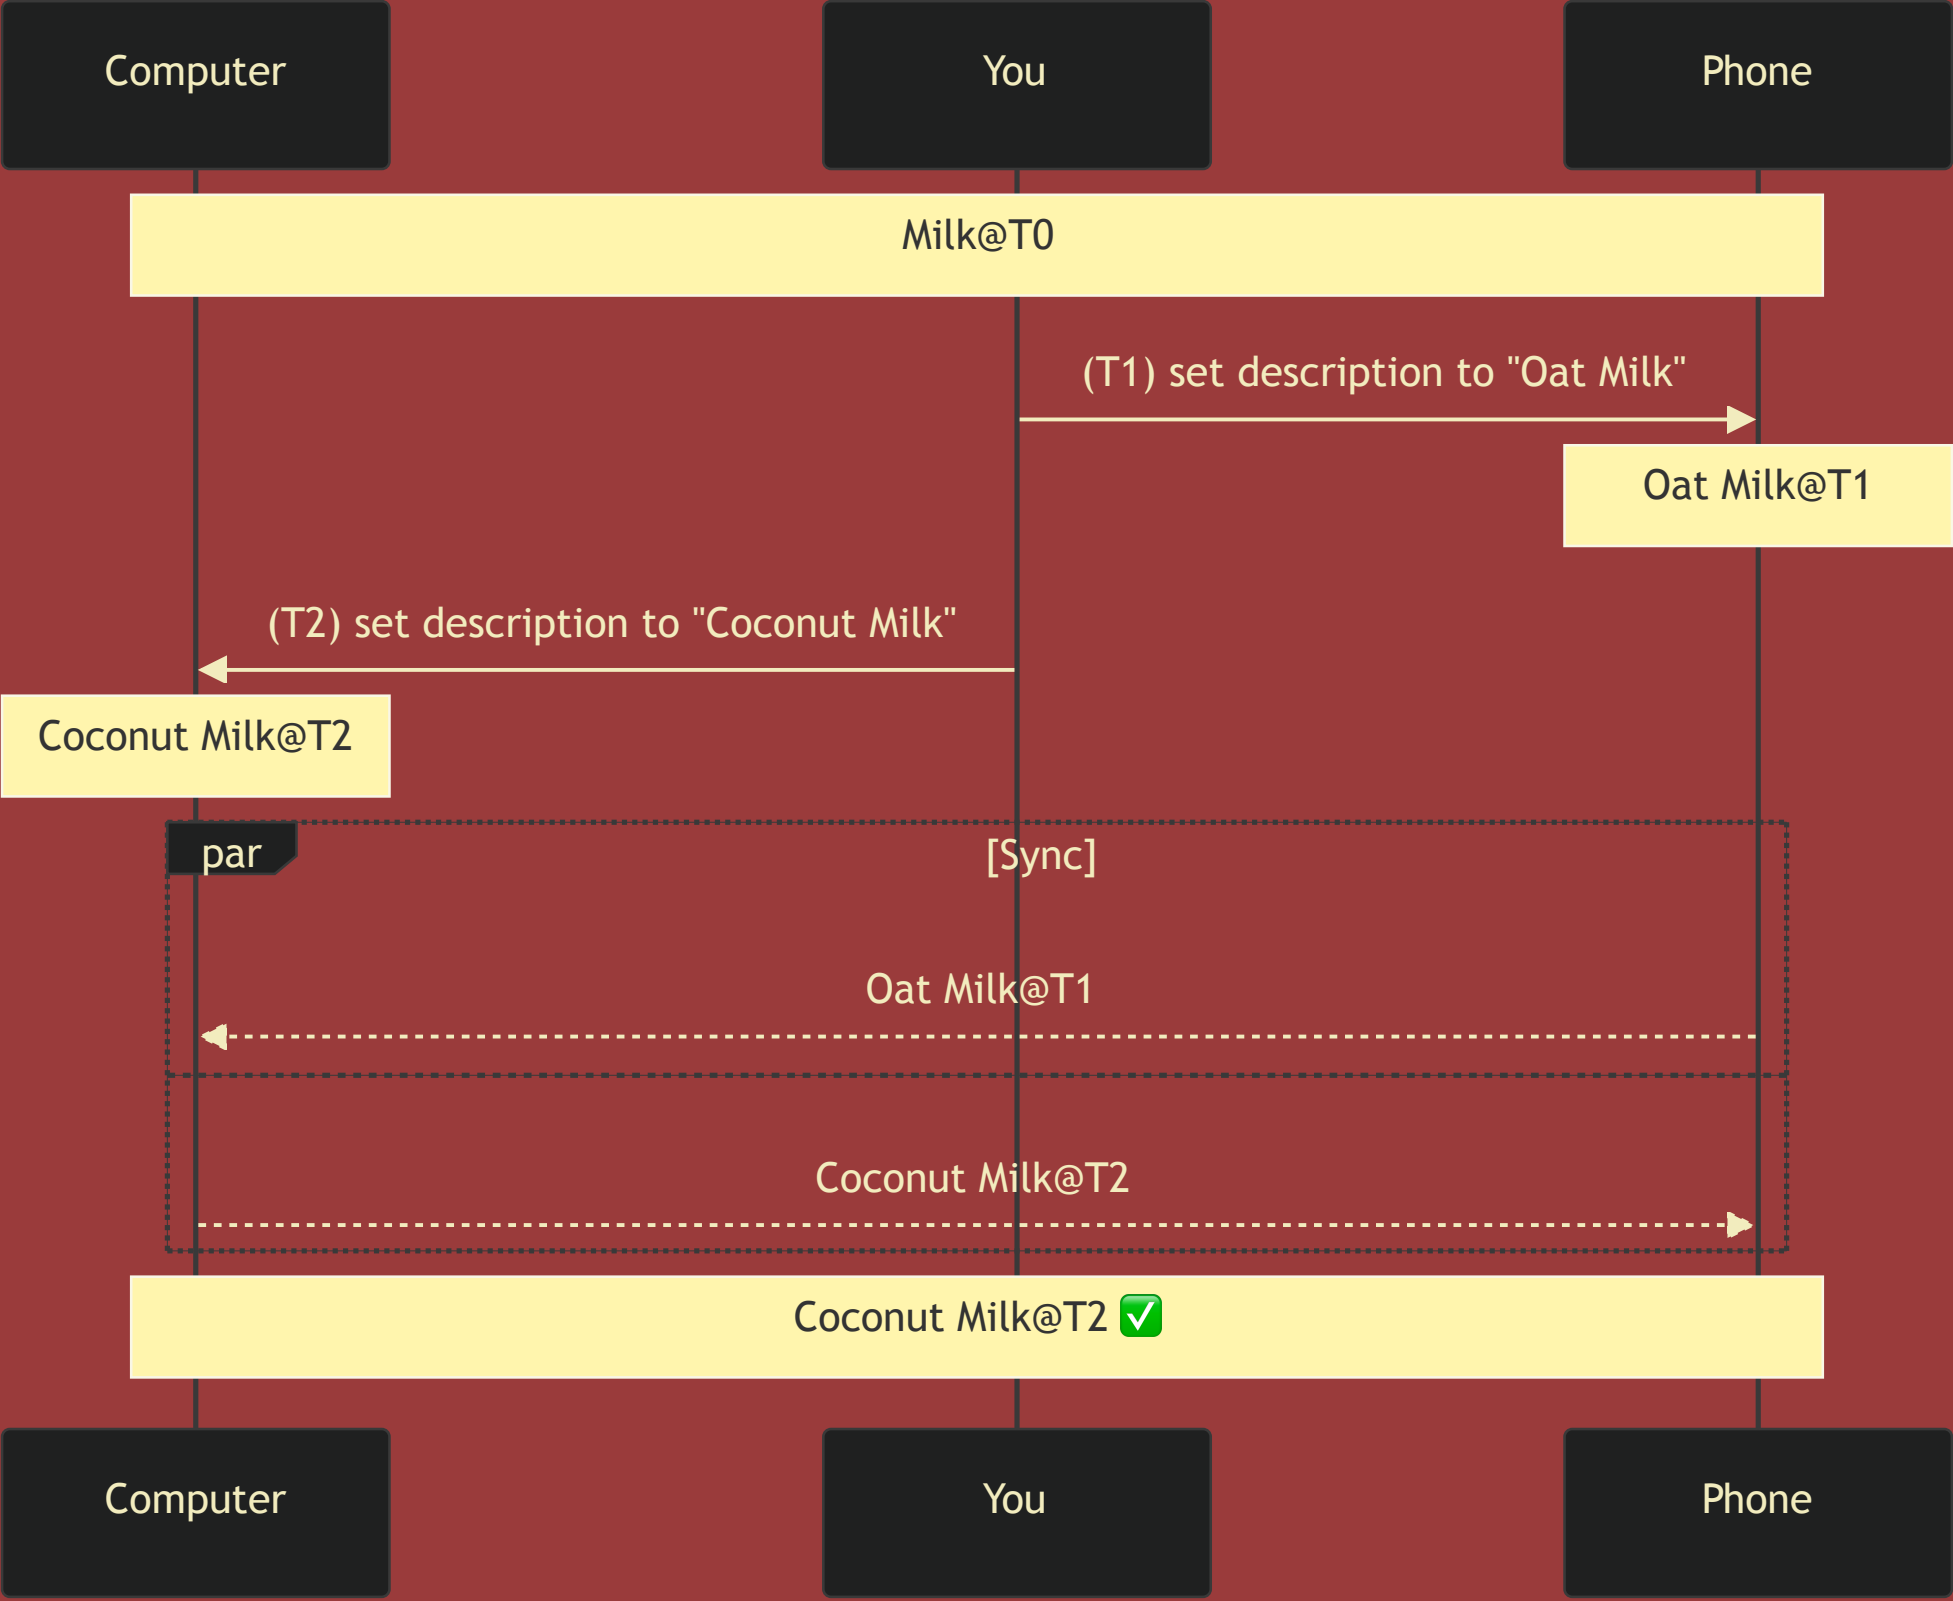
```
pub struct HybridLogicalClock {  
    timestamp: DateTime<Utc>,  
    counter: u16,  
    node_id: Uuid,  
}  
  
impl HybridLogicalClock {  
    pub fn tick(&mut self) {  
        let now = Utc::now();  
        if now > self.timestamp {  
            self.timestamp = now;  
            self.counter = 0;  
        } else {  
            self.counter += 1;  
        }  
    }  
}
```

Hybrid Logical Clock

```
impl Ord for HybridLogicalClock {  
    fn cmp(&self, other: &Self) -> Ordering {  
        self.timestamp.cmp(&other.timestamp)  
            .then(self.counter.cmp(&other.counter))  
            .then(self.node_id.cmp(&other.node_id))  
    }  
}
```


LWW Register

```
pub struct LWWRegister<T> {  
    value: T,  
    clock: HybridLogicalClock,  
}  
  
impl<T> Merge for LWWRegister<T> {  
    fn merge_mut(&mut self, other: Self) {  
        if other.clock > self.clock {  
            self.value = other.value;  
            self.clock = other.clock;  
        }  
    }  
}
```



Hold on, wasn't that a conflict?

- Yeah, semantically.
- CRDTs only guarantee that values converge, not that they make sense.
- One way to solve this: multi-value register that allows the user to pick what they wanted.
- Another way: use a sequence CRDT to edit text. Resolve to something like "Coconut Oat Milk"
- Choosing the right CRDTs for your application is about choosing the constraints you can live with.

Grow-Only Map

```
pub struct GMap<K: Hash + Ord, V: Merge>(BTreeMap<K, V>);
```

Grow-Only Map

```
impl<K: Hash + Ord, V: Merge> GMap<K, V> {  
    pub fn insert(&mut self, key: K, value: V) {  
        match self.0.entry(key) {  
            Entry::Occupied(mut existing) => {  
                existing.get_mut().merge_mut(value);  
            }  
            Entry::Vacant(vacant) => {  
                vacant.insert(value);  
            }  
        }  
    }  
}
```

Grow-Only Map

```
impl<K: Hash + Ord, V: Merge> Merge for GMap<K, V> {  
    fn merge_mut(&mut self, other: Self) {  
        for (key, value) in other.0 {  
            self.insert(key, value);  
        }  
    }  
}
```

**We now have
enough to make
our to-do list!**

Document, V1

```
pub struct Document {  
    tasks: GMap<Uuid, Task>,  
}  
  
pub struct Task {  
    added: LWWRegister<DateTime<Utc>>,  
    description: LWWRegister<String>,  
    completed: LWWRegister<bool>,  
}
```


Document, V1

```
impl Merge for Document {  
    fn merge_mut(&mut self, other: Self) {  
        self.tasks.merge_mut(other.tasks);  
    }  
}  
  
impl Merge for Task {  
    fn merge_mut(&mut self, other: Self) {  
        self.added.merge_mut(other.added);  
        self.complete.merge_mut(other.complete);  
        self.description.merge_mut(other.description);  
    }  
}
```

Speaking of
constraints...

How about removing values?

A "two-phase map", or 2P-Map

```
pub struct TwoPMap<K: Ord + Clone, V: Merge> {  
    adds: BTreeMap<K, V>,  
    removes: BTreeSet<K>,  
}
```

How about removing values?

```
impl<K: Ord + Clone, V: Merge> TwoPMap<K, V> {  
    pub fn insert(&mut self, key: K, value: V) {  
        if self.removes.contains(&key) {  
            return;  
        }  
  
        // (snip) same code as GMap::insert  
    }  
  
    pub fn remove(&mut self, key: K) {  
        self.adds.remove(&key);  
        self.removes.insert(key);  
    }  
}
```

How about removing values?

```
impl<K: Ord + Clone, V: Merge> Merge for TwoPMap<K, V> {  
    fn merge_mut(&mut self, mut other: Self) {  
        // same as GSet  
        self.removes.append(&mut other.removes);  
  
        // same as GMap  
        for (key, value) in other.adds {  
            self.insert(key, value);  
        }  
  
        // drop unnecessary values  
        self.adds.retain(|k, _| !self.removes.contains(k))  
    }  
}
```

Metadata Overhead

In general, this is a problem.

- 2P-Sets and 2P-Maps keep tombstones forever.
- G-Sets and G-Maps can never shrink.
- LWW Registers keep an HLC (ours are 34 bytes each.)
- Sequence CRDTs have a per-item overhead. (So "Hello" would have 5 characters worth)
- Compression helps.
- Optimizations in industrial CRDTs (e.g. ORSWOT) mostly go towards avoiding this problem.

Archiving Completed Tasks

```
pub struct Document {  
-   tasks: GMap<Uuid, Task>,  
+   tasks: TwoPMap<Uuid, Task>,  
}
```

Bookkeeping

```
pub struct Replica {  
    id: Uuid,  
    clock: HybridLogicalClock,  
    document: Document,  
}  
  
impl Replica {  
    pub fn receive(&mut self, other: Replica) {  
        self.document.merge_mut(other.document);  
        self.clock = self.clock.max(other.clock).claim(self.id);  
    }  
}
```


Demo !

Thanks



- bytes.zone
- github.com/BrianHicks/rust-crdt-talk

Implementations:

- automerge.org
- loro.dev
- crates.io/crates/crdts